



ECCE 632: Advanced Operating Systems
Assignment

System Calls

Group Members:

Ghebrebrhan Weldit	100043383
Zhang Zhibo	100060990
Mohamed Tarek	100038660

Submitted to:

Dr Ernesto Damiani
Spring 2022

1	INTRODUCTION	3
1.1	System Calls	3
2	PREPARE LINUX KERNEL	5
2.1	Preparation	5
2.2	Configuration	7
3	ADDING A HELLO WORLD SYS CALL	8
3.1	Compiling	9
4	SYSTEM CALL – PROCSCHED	11
4.1	Kernel Module	12
4.2	Compiling (The same as in adding a helloworld)	13
4.3	Test MSSV	14
5	WRAPPER	16
5.1	Create “procsched.h” and “procsched.c” files.	16

1 Introduction

1.1 System Calls

System calls provide an interface for user programs to access the services made available by the operating system. Application developers access system calls via an application programming interface (API). The API specifies a set of C functions that are available to the application programmer, including the parameters that are passed to each function and the return values the programmer can expect. The functions that make up the API typically invoke the actual system calls on behalf of the application programmer.

1.1.1 Example of Standard API

As an example of the standard POSIX API, consider the `open()` function that is available in UNIX and Linux systems. The API for this function is obtained from the man page. A description of this API appears below:

```
$ man open

#include <unistd.h>

int open(const char *path, int oflags);
int open(const char *path, int oflags, mode_t mode);
```

A program that uses the `open()` function must include the `unistd.h` header file, as this file defines `int` data types (among other things). The parameters passed to `open()` are as follows.

- `const *path` - The relative or absolute path to the file that is to be opened.
- `int oflags` - A bitwise 'or' separated list of values that determine the method in which the file is to be opened.
- `mode_t mode` - A bitwise 'or' separated list of values that determine the permissions of the file if it is created.

The file descriptor returned is the integer that is positive or zero if the file was opened without problems. If a negative value is returned, then there was an error occur while opening the file.




Figure 1: A user application invoking the `open()` system call.

The relationship between an API, the system-call interface, and the operating system is shown in Figure 1, which illustrates how the operating system handles a user application invoking the `open()` system call.

1.1.2 Diagram

Figure 2 shows, the diagram of doing the assignment. First, you will practice the progress of compiling Linux kernel. After that, the most important part is to implement a system call inside the kernel. The assignment is divided into multiple stages and score is marked by each stage as Figure 2.

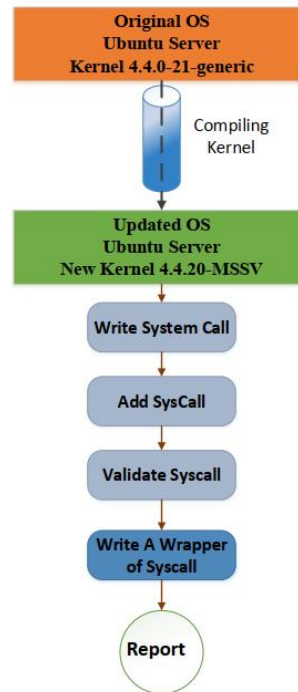


Figure 2: Diagram of implementing the assignment

2 Prepare Linux Kernel

The main objective of this chapter is to prepare the linux kernel on a virtual machine running Ubuntu 20.04.

2.1 Preparation

2.1.1 Set up Virtual Machine

1. Oracle VM Virtual Box Manager was used along with Ubuntu 20.4 image.

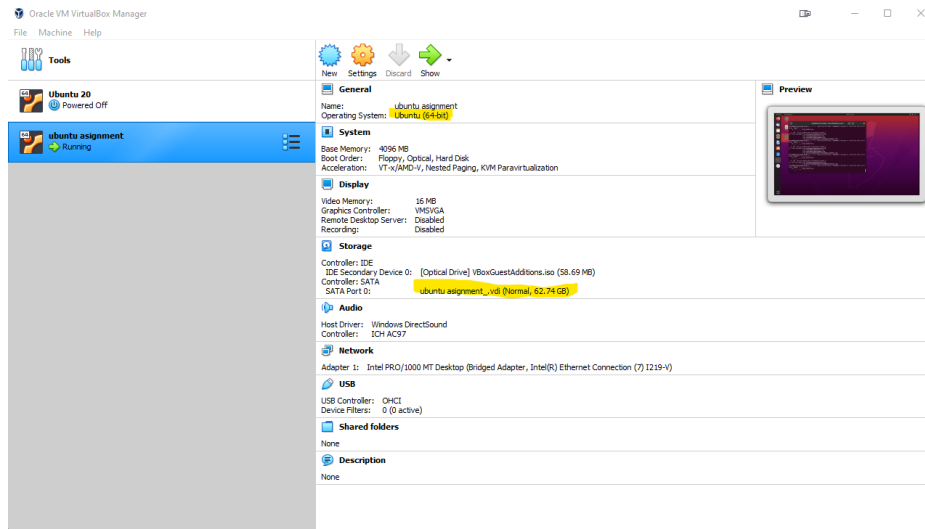


Figure 3: Virtual Box Manager Settings

Some settings were modified to allow us the system to run smoothly. First, at least 50 GB of storage memory is required to be allocated for Ubuntu to allow all the required packages to be installed. Secondly, bidirectional drag and drop and clipboarding to allow the host machine to send/share data to the virtual machine. For networking, it is important to set bridge adapter to allow the virtual machine to access the internet to download the required packages.

2.1.2 Install core packages

- Get Ubuntu's toolchain (gcc, make, and so forth) by installing the **build essential** metapackage:

2.1.3

```
$ sudo apt-get update
$ sudo apt-get install build-essential
```

- Install kernel-package:

```
$ sudo apt-get install kernel-package
```

2.1.4 Create a kernel compilation directory:

It is recommended to create a separate build directory for your kernel(s). In this example,

the directory kernelbuild will be created in the home directory:

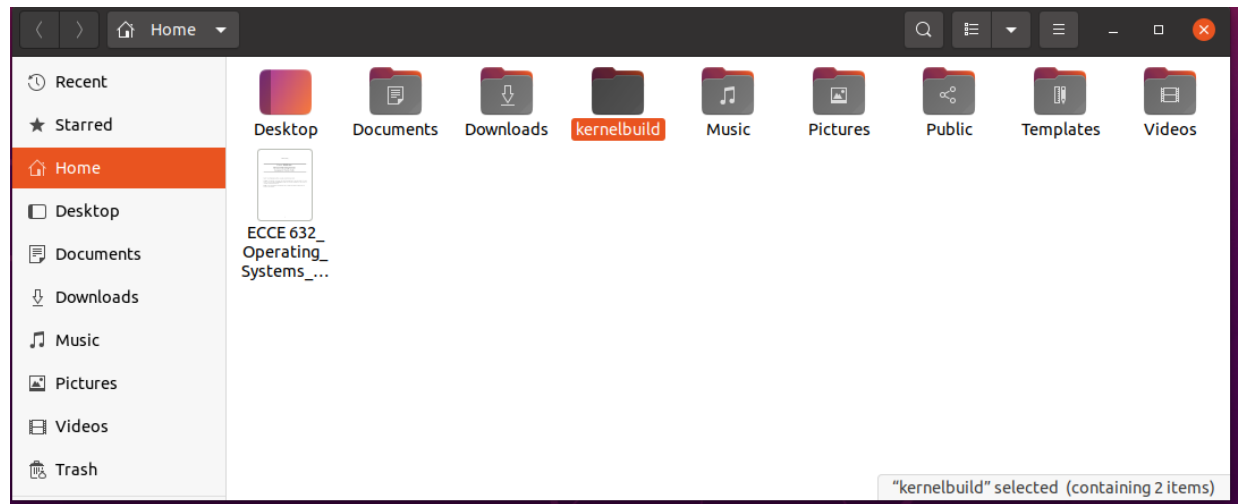


Figure 4: Kernel Directory

2.1.5 Download the kernel source

Download the kernel source from <http://www.kernel.org>. This should be the tarball (tar.xz) file for your chosen kernel. It can be downloaded by simply right-clicking the tar.xz link in your browser and selecting Save Link As..., or any other number of ways via alternative graphical or command-line tools that utilize HTTP, FTP, RSYNC, or Git. In the following command-line example, wget has been installed and is used inside the `~/kernelbuild` directory to obtain kernel 4.4.56.

```
$ cd ~/kernelbuild
$ wget https://cdn.kernel.org/pub/linux/kernel/v4.x/linux-4.4.56.tar.xz
```

QUESTION: Why we have to use another kernel source from the server such as <http://www.kernel.org>, can we compile the original kernel (the local kernel on the running OS) directly?

We can't compile the local operating system because its already compiled. The local operating system is in a binary format and couldn't be compiled. I need to download another copy, source format, then I need to build the binary of the source code in each directory on my disk.

2.2 Configuration

A crucial step in customizing the default kernel to reflect your computer's precise specifications. Kernel configuration is its .config file, which includes the use of Kernel modules. By setting the options in .config properly, your kernel and computer will function most efficiently. Since making our own configuration file is a complicated process, we could borrow the content of configuration file of an existing kernel currently used by the virtual machine. This file is typically located in /boot/ so our job is simply copy it to the source code directory:

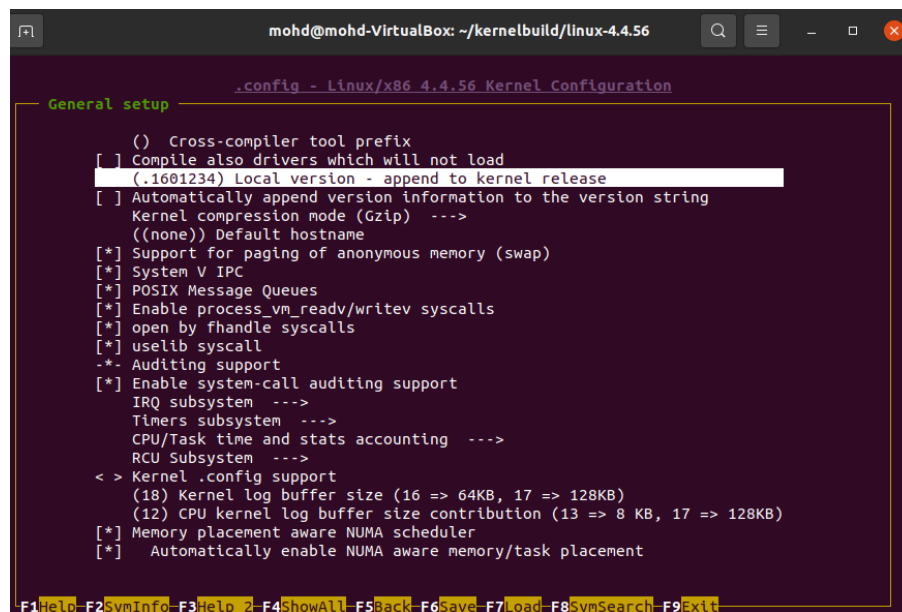
```
$ cp /boot/config-x.x.x-generic ~/kernelbuild/linux-4.4.56/.config
```

Using the command `uname -r` we can find the specific version of kernel installed.

```
mohd@mohd-VirtualBox:~/kernelbuild/linux-4.4.56$ uname -r
5.13.0-44-generic
```

run `$ make menuconfig` or `$ make nconfig` inside the top directory to open Kernel Configuration.

To change kernel version, go to General setup option, Access to the line “(-ARCH) Local version append to kernel release”. Then enter a dot “.” followed by your MSSV. For example: .1601234 Press F6 to save your change and then press F9 to exit.



Install relevant compilers:

```
sudo apt-get update
sudo apt-get install libncurses5-dev libssl-dev
sudo apt-get install build-essential openssl
sudo apt-get install zlibc minizip
sudo apt-get install bison
sudo apt-get install vim
sudo apt-get install flex
sudo apt-get install libssl-dev
sudo apt install dwarves
sudo apt-get install libelf-dev
```

3 Adding a Hello World Sys call

Add the desired system call function

```
sudo vim sys.c
```

(Note: In vim, “I” enters edit, “esc” exits edit. “G” jumps to the end, “gg” goes to the beginning. “:wq” save to exit, “:q” don’t save to exit)

Add the following function in this file:

```
asm linkage long sys_helloworld(void)
{
    printk("helloworld!");
    return 1;
}
```

Add the system call number:

```
439 64 helloworld sys_helloworld
```

```
cd /usr/src/linux-5.6
sudo make mrproper
sudo make clean
```

```
sudo make menuconfig
```

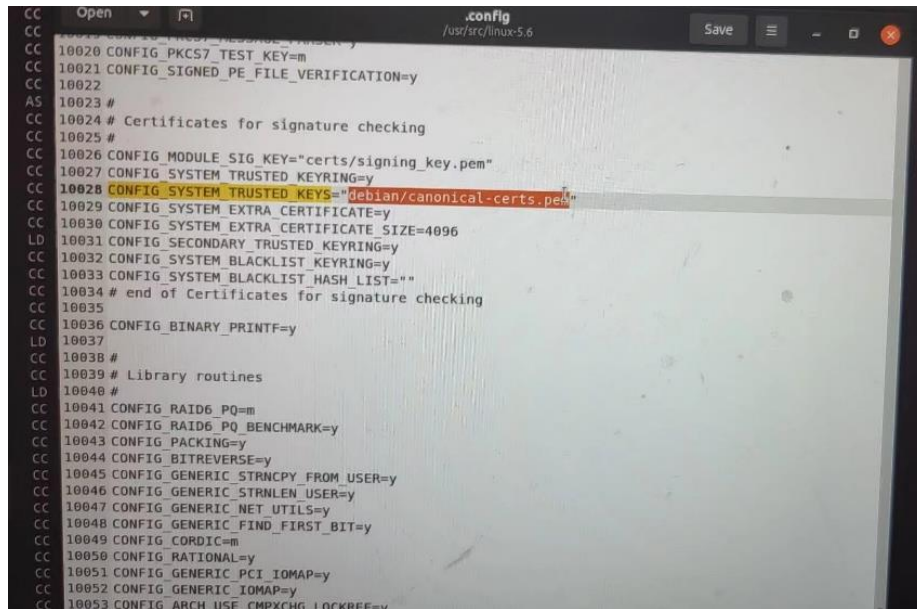
since there is a runtime error, we had to troubleshoot by using `sudo gedit .config`

Find:

```
CONFIG_SYSTEM_TRUSTED_KEYS="debian/canonical-certs.pem"
```

Delete:

```
debian/canonical-certs.pem
```

```
CC Open  .config  /usr/src/linux-5.6  Save
CC 10020 CONFIG_PKCS7_TEST_KEY=m
CC 10021 CONFIG_SIGNED_PE_FILE_VERIFICATION=y
CC 10022
AS 10023 #
CC 10024 # Certificates for signature checking
CC 10025 #
CC 10026 CONFIG_MODULE_SIG_KEY="certs/signing_key.pem"
CC 10027 CONFIG_SYSTEM_TRUSTED_KEYRING=y
CC 10028 CONFIG_SYSTEM_TRUSTED_KEYS="debian/canonical-certs.pem"
CC 10029 CONFIG_SYSTEM_EXTRA_CERTIFICATE=y
CC 10030 CONFIG_SYSTEM_EXTRA_CERTIFICATE_SIZE=4096
CC 10031 CONFIG_SECONDARY_TRUSTED_KEYRING=y
CC 10032 CONFIG_SYSTEM_BLACKLIST_KEYRING=y
CC 10033 CONFIG_SYSTEM_BLACKLIST_HASH_LIST=""
CC 10034 # end of Certificates for signature checking
CC 10035
CC 10036 CONFIG_BINARY_PRINTF=y
LD 10037
CC 10038 #
CC 10039 # Library routines
LD 10040 #
CC 10041 CONFIG_RAID6_PQ=m
CC 10042 CONFIG_RAID6_PQ_BENCHMARK=y
CC 10043 CONFIG_PACKING=y
CC 10044 CONFIG_BITREVERSE=y
CC 10045 CONFIG_GENERIC_STRNCPY_FROM_USER=y
CC 10046 CONFIG_GENERIC_STRNLEN_USER=y
CC 10047 CONFIG_GENERIC_NET_UTILS=y
CC 10048 CONFIG_GENERIC_FIND_FIRST_BIT=y
CC 10049 CONFIG_CORDIC=m
CC 10050 CONFIG_RATIONAL=y
CC 10051 CONFIG_GENERIC_PCI_IOMAP=y
CC 10052 CONFIG_GENERIC_IOMAP=y
CC 10053 CONFIG_ARCH_USE_CMPXCHG_LOCKREF=y
```

3.1 Compiling

sudo make -j4 (to use 4 cores)

(The terminal window is maximized when performing the steps in this block, otherwise an error may be reported)

It will take 2-3 hours.


sudo make modules_install

sudo make install

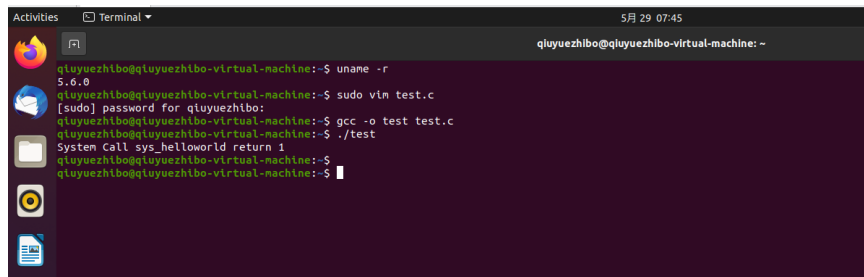
sudo reboot

(After rebooting, click the mouse to enter ubuntu and quickly hold down shift and long press)

Choose the linux version.

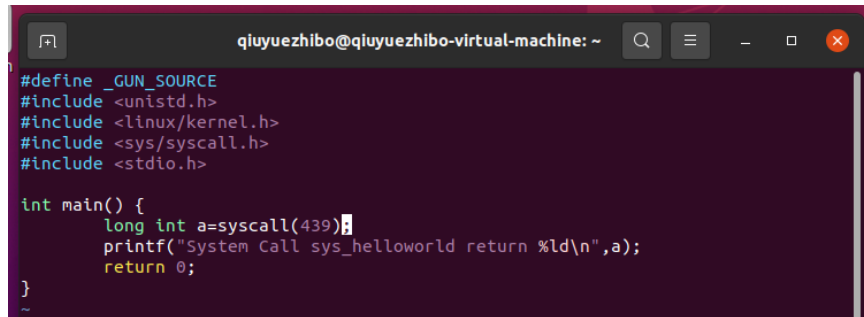


Test by opening the terminal and running command **uname -r**



A terminal window titled 'Terminal' with a date and time of '5月 29 07:45'. The prompt is 'qluyuezhibo@qluyuezhibo-virtual-machine: ~'. The user enters 'uname -r' and the output is '5.6.0'. Then, the user enters 'sudo vim test.c', followed by a password prompt '[sudo] password for qluyuezhibo:'. The user then enters 'gcc -o test test.c', followed by './test'. The output of the program is 'System Call sys_helloworld return 1'. The prompt returns to '\$'.

```
qluyuezhibo@qluyuezhibo-virtual-machine:~$ uname -r
5.6.0
qluyuezhibo@qluyuezhibo-virtual-machine:~$ sudo vim test.c
[sudo] password for qluyuezhibo:
qluyuezhibo@qluyuezhibo-virtual-machine:~$ gcc -o test test.c
qluyuezhibo@qluyuezhibo-virtual-machine:~$ ./test
System Call sys_helloworld return 1
qluyuezhibo@qluyuezhibo-virtual-machine:~$
```

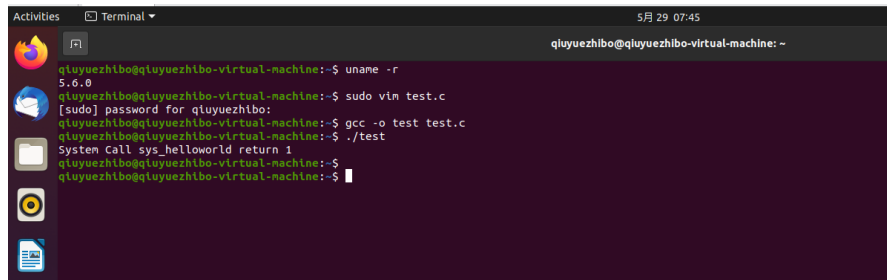


A code editor window titled 'qluyuezhibo@qluyuezhibo-virtual-machine: ~' with a search bar and window controls. It displays the source code of 'test.c'.

```
#define _GNU_SOURCE
#include <unistd.h>
#include <linux/kernel.h>
#include <sys/syscall.h>
#include <stdio.h>

int main() {
    long int a=syscall(439);
    printf("System Call sys_helloworld return %ld\n",a);
    return 0;
}
```

Run the command **gcc -o test test.c**



A terminal window titled 'Terminal' with a date and time of '5月 29 07:45'. The prompt is 'qluyuezhibo@qluyuezhibo-virtual-machine: ~'. The user enters 'uname -r' and the output is '5.6.0'. Then, the user enters 'sudo vim test.c', followed by a password prompt '[sudo] password for qluyuezhibo:'. The user then enters 'gcc -o test test.c', followed by './test'. The output of the program is 'System Call sys_helloworld return 1'. The prompt returns to '\$'.

```
qluyuezhibo@qluyuezhibo-virtual-machine:~$ uname -r
5.6.0
qluyuezhibo@qluyuezhibo-virtual-machine:~$ sudo vim test.c
[sudo] password for qluyuezhibo:
qluyuezhibo@qluyuezhibo-virtual-machine:~$ gcc -o test test.c
qluyuezhibo@qluyuezhibo-virtual-machine:~$ ./test
System Call sys_helloworld return 1
qluyuezhibo@qluyuezhibo-virtual-machine:~$
```

Success.

4 System Call – procsched

In this assignment, we have to define a system call named ‘procsched’. This syscall helps user to show the schedule information of a specific process. An example output:

```
pcount: 4567
run_delay: 829035007
last_arrival: 1818402352768
last_queued: 0
```

To implement this system call, you have to use 2 data structures defined by Linux OS, task struct and sched info. In the Linux kernel, every process has an associated struct task_struct. The definition of this struct is in the header file include /linux/sched.h

```
struct task_struct {
    volatile long state;
    /* -1 unrunnable, 0 runnable, >0 stopped */
    struct thread_info *thread_info;
    atomic_t usage;
    ...
    ...
    struct sched_info sched_info;
    ...
    char comm[16];
    ...
};
```

The sched_info within the task_struct describe schedule information of the process. The sched_info is defined in include /linux/sched.h as:

```
struct sched_info {
#ifdef CONFIG_SCHED_INFO
    /* # of times we have run on this CPU: */
    unsigned long pcount;
    /* Time spent waiting on a runqueue: */
    unsigned long long run_delay;
    /* When did we last run on a CPU? */
    unsigned long long last_arrival;
    /* When were we last queued to run? */
    unsigned long long last_queued;
#endif /* CONFIG_SCHED_INFO */
};
```

Kernel modules must have at least two functions: a “start” (initialization) function called init module() which is called when the module is insmoded into the kernel, and an “end” (cleanup) function called cleanup module() which is called just before it is rmmoded.

QUESTION: What is the meaning of other parts, i.e., i386, Procsched, and sys Procsched?

I386: ABI (Application Binary Interface)

Procsched: none of the syscall

Sys_procsched: entry point.

4.1 Kernel Module

Add the desired system call function

`sudo vim sys.c`

Add following codes: (which are basically the codes in “sys_procsched.c” file.

```
#include <linux/linkage.h>
#include <linux/sched.h>
#include <linux/kernel.h>
#include <linux/uaccess.h>

struct proc_segs {
    unsigned long mssv;
    unsigned long pcount;
    unsigned long long run_delay;
    unsigned long long last_arrival;
    unsigned long long last_queued;
};

asmlinkage long sys_procsched(int pid, struct proc_segs * info){
    struct task_struct * task;
    printk("Finding...\n");
    for_each_process(task) {
        printk("[%d] ----- [%s]\n", task->pid, task->comm);
        if(task->pid == pid) {
            if(task->mm != NULL) {
                struct proc_segs buff;
                printk("inside process!\n");
                buff.mssv = 1601234;
                buff.pcount = task->sched_info.pcount;
                buff.run_delay = task->mm->run_delay;
                buff.last_arrival = task->mm->last_arrival;
                buff.last_queued = task->mm->last_queued;
                int res = copy_to_user(info, &buff, sizeof(buff));
                if(res == 0) printk("copy data successful!\n");
                else printk("copy data failed\n");
                printk("Find out pid [%d]", pid);
                return 0;
            }
        }
    }
    return -1;
}
```

4.1.1 Add the statement

`cd /usr/src/linux-5.6/arch/x86/include/asm/`
`sudo vim syscalls.h`

```

#include <linux/compiler.h>
#include <linux/linkage.h>
#include <linux/signal.h>
#include <linux/types.h>

/* Common in X86_32 and X86_64 */
/* kernel/loport.c */
long ksys_loperm(unsigned long from, unsigned long mun, int turn_on);

#ifdef CONFIG_X86_32
/*
 * These definitions are only valid on pure 32-bit systems; x86-64 uses a
 * different syscall calling convention
 */
asmlinkage long sys_loperm(unsigned long, unsigned long, int);
asmlinkage long sys_lopl(unsigned int);
asmlinkage long sys_helloworld(void);
struct proc_segs;
asmlinkage long sys_procsched(int pid, struct proc_segs * info);
/* kernel/ldt.c */
asmlinkage long sys_modify_ldt(int, void __user *, unsigned long);
/* kernel/signal.c */
asmlinkage long sys_rt_sigreturn(void);

/* kernel/tls.c */
asmlinkage long sys_set_thread_area(struct user_desc __user *);
asmlinkage long sys_get_thread_area(struct user_desc __user *);

/* X86_32 only */
...
/* kernel/signal.c */
asmlinkage long sys_sigreturn(void);
-- INSERT --

```

Add:

struct proc_segs;

asmlinkage long sys__procsched(int pid, struct proc_segs * info)

4.1.2 Add system call number

cd /usr/src/linux-5.6/arch/x86/entry/syscalls

sudo vim syscall_64.tbl

519	x32	recvmsg	__x32_compat_sys_recvmsg
520	x32	execve	__x32_compat_sys_execve/ptregs
521	x32	ptrace	__x32_compat_sys_ptrace
522	x32	rt_sigpending	__x32_compat_sys_rt_sigpending
523	x32	rt_sigtimedwait	__x32_compat_sys_rt_sigtimedwait_time64
524	x32	rt_sigqueueinfo	__x32_compat_sys_rt_sigqueueinfo
525	x32	sigaltstack	__x32_compat_sys_sigaltstack
526	x32	timer_create	__x32_compat_sys_timer_create
527	x32	mq_notify	__x32_compat_sys_mq_notify
528	x32	kexec_load	__x32_compat_sys_kexec_load
529	x32	waitid	__x32_compat_sys_waitid
530	x32	set_robust_list	__x32_compat_sys_set_robust_list
531	x32	get_robust_list	__x32_compat_sys_get_robust_list
532	x32	vmsplice	__x32_compat_sys_vmsplice
533	x32	move_pages	__x32_compat_sys_move_pages
534	x32	preadv	__x32_compat_sys_preadv64
535	x32	pwritev	__x32_compat_sys_pwritev64
536	x32	rt_tgsigqueueinfo	__x32_compat_sys_rt_tgsigqueueinfo
537	x32	recvmsg	__x32_compat_sys_recvmsg_time64
538	x32	sendmsg	__x32_compat_sys_sendmsg
539	x32	process_vm_readv	__x32_compat_sys_process_vm_readv
540	x32	process_vm_writev	__x32_compat_sys_process_vm_writev
541	x32	setsockopt	__x32_compat_sys_setsockopt
542	x32	getsockopt	__x32_compat_sys_getsockopt
543	x32	io_setup	__x32_compat_sys_io_setup
544	x32	io_submit	__x32_compat_sys_io_submit
545	x32	execveat	__x32_compat_sys_execveat/ptregs
546	x32	preadv2	__x32_compat_sys_preadv64v2
547	x32	pwritev2	__x32_compat_sys_pwritev64v2
548	x32	procsched	__sys_procsched

548 x32 procsched sys_procsched

4.2 Compiling (The same as in adding a helloworld)

cd /usr/src/linux-5.6

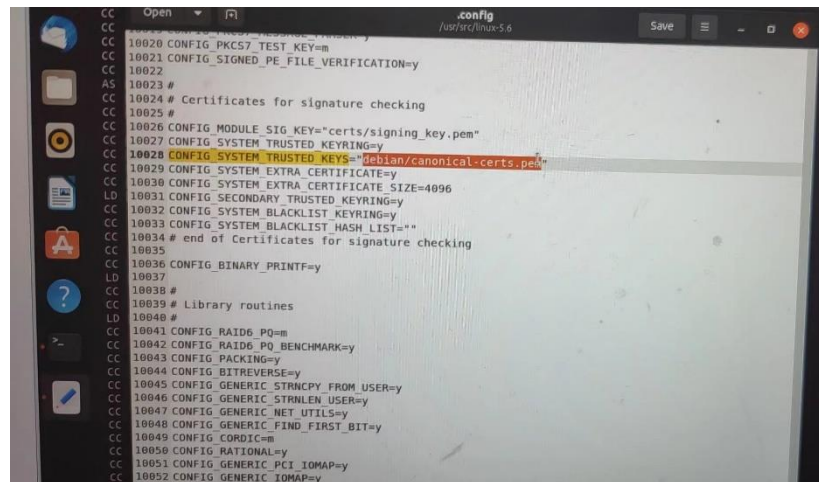
sudo make mrproper

sudo make clean

sudo make menuconfig

When the window appears use the left and right keyboard keys to control directly save->ok->exit->exit

`sudo gedit .config`



Find:

`CONFIG_SYSTEM_TRUSTED_KEYS=\"debian/canonical-certs.pem\"`

Delete:

`debian/canonical-certs.pem`

`sudo make -j4`

(The terminal window is maximised when performing the steps in this block, otherwise an error may be reported)

It will take 2-3 hours.

`sudo make modules_install`

`sudo make install`

`sudo reboot`

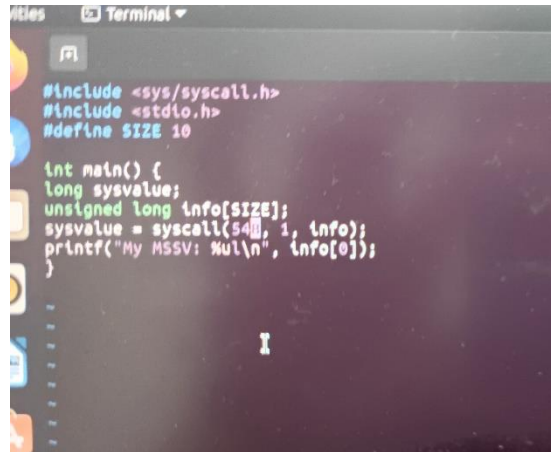
(After rebooting, click the mouse to enter ubuntu and quickly hold down shift and long press)

4.3 Test MSSV

`sudo vim test2.c`

Add following codes:

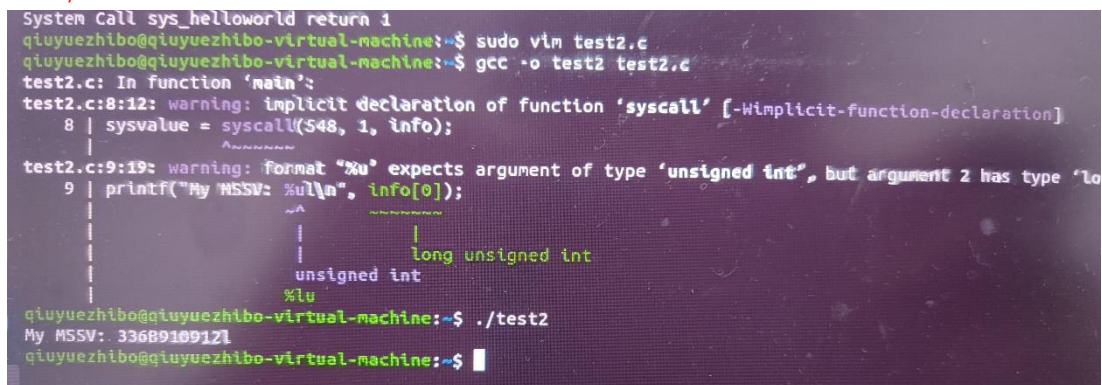
```
#include <sys/syscall.h>
#include <stdio.h>
#define SIZE 10
int main() {
    long sysvalue;
    unsigned long info[SIZE];
    sysvalue = syscall(548, 1, info);
    printf("My MSSV: %ul\n", info[0]);
}
```

```
#include <sys/syscall.h>
#include <stdio.h>
#define SIZE 10

int main() {
    long sysvalue;
    unsigned long info[SIZE];
    sysvalue = syscall(548, 1, info);
    printf("My MSSV: %u\n", info[0]);
}
```

gcc -o test2 test2.c
./test2



```
System Call sys_helloworld return 1
qiuyuezhibo@qiuyuezhibo-virtual-machine:~$ sudo vim test2.c
qiuyuezhibo@qiuyuezhibo-virtual-machine:~$ gcc -o test2 test2.c
test2.c: In function 'main':
test2.c:8:12: warning: implicit declaration of function 'syscall' [-Wimplicit-function-declaration]
8 | sysvalue = syscall(548, 1, info);
  |             ^~~~~~
test2.c:9:19: warning: format '%u' expects argument of type 'unsigned int', but argument 2 has type 'long unsigned int' [-Wformat]
9 | printf("My MSSV: %u\n", info[0]);
  |                   ^~
  |                   |
  |                   long unsigned int
  |                   unsigned int
  |                   %lu
qiuyuezhibo@qiuyuezhibo-virtual-machine:~$ ./test2
My MSSV: 33689109121
qiuyuezhibo@qiuyuezhibo-virtual-machine:~$
```

Succeed!
MSSV shown on the screen.

Open the file include/linux/syscalls.h and add the following line to the end of this file:

```
struct proc_segs;

asmlinkage long sys_procsched(int pid, struct proc_segs * info);
```

QUESTION: What is the meaning of each line above?

The 1st Line: struct proc_segs:

Proc_segs argument of the type struct

The 2nd line:

Long: function handling syscall return

Asmlinkage is a #define with some gcc magic tells the compiler that the function is not expected to optimally find all the arguments in the registers (a common optimization), but only on the CPU's stack

System_call consume its first argument, the system call number and allows up to 4 more arguments passed to real system, on the stack, all system calls are marked with asmlinkage tag, so they all look to the stack for arguments.

Its also used to allow calling a function from assembly files.

First run “make” to compile the kernel and create vmlinuz. It takes a long time to “\$ make”, we can run this stage in parallel by using tag “-j np”, where np is the number of processes you run this command.

```
$ make
or
$ make -j 4
```

vmlinuz is “the kernel”. Specifically, it is the kernel image that will be uncompressed and loaded into memory by GRUB or whatever other boot loader you use.

Then build the loadable kernel modules. Similarly, you can run this command in parallel.

```
$ make modules
or
$ make -j 4 modules
```

QUESTION: What is the meaning of these two stages, namely “make” and “make modules”?

“Make”: to compile and link the kernel image, there is a single file name “vmlinuz”

“Make modules” compiles individual files for each question you answered “M” during kernel config. The object code is linked against your freshly built kernel. For each question answered “Y” there are already part of “vmlinuz” and questions answered “N” are skipped.

After compiling and executing this program, your MSSV should be shown on the screen.

QUESTION: Why this program could indicate whether our system works or not?

The program is to check if system call has been integrated into kernel by calling syscall ([number 32], 1, Info, info[0] should be the MSSV. Because MSSV is defined in the first line of the struct proc_segs

5 Wrapper

5.1 Create “procsched.h” and “procsched.c” files.

procsched.c:

```
#include "procsched.h"
#include <linux/kernel.h>
#include <sys/syscall.h>
#define _SYS_PROCMEM 548

long procsched(pid_t pid, struct proc_segs * info) {
    long sysvalue;
    sysvalue = syscall(_SYS_PROCMEM, pid, info);
    return sysvalue;
}
```


Note: You must define fields in `proc_segs` struct in the same order as you did in the kernel.

QUESTION: Why we have to re-define `proc_segs` struct while we have already defined it inside the kernel?

Because we have to leave out kernel source code directory and create to store the source for our wrapper. The redefine of `proc_segs` is for the prototype of the wrapper. They are in different directory with the one defined inside the kernel.

`procsched.h:`

```
#ifndef _PROC_MEM_H_
#define _PROC_MEM_H_
#include <unistd.h>

struct proc_segs {
    unsigned long mssv;
    unsigned long pcount;
    unsigned long long run_delay;
    unsigned long long last_arrival;
    unsigned long long last_queued;
};

long procsched(pid_t pid, struct proc_segs * info);

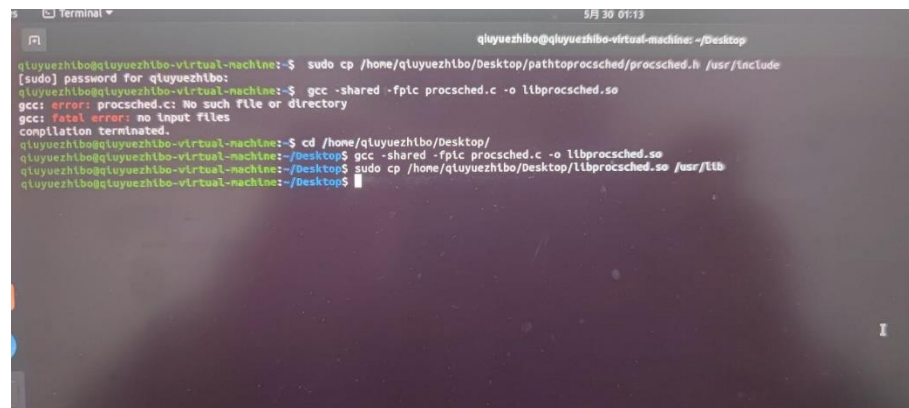
#endif // _PROC_SCHED_H_
```

Validation You could check your work by write an additional test module to call this functions but do not include the test part to your source file (`procsched.c`). After making sure that the wrapper work properly, we then install it to our virtual machine. First, we must ensure everyone could access this function by making the header file visible to GCC. Run following command to copy our header file to header directory of our system:

```
$ sudo cp <path to procsched.h> /usr/include
```

QUESTION: Why root privilege (e.g. adding `sudo` before the `cp` command) is required to copy the header file to `/usr/include`?

The command “`CD`” needs to be run with root privilege. We are copying our header file “`procsched.h`” top header directory of our system so that everyone could access this function. This changed the privilege of the file, therefore, needs the root privilege “`sudo`”



```
qiyuezhibo@qiyuezhibo-virtual-machine: ~/Desktop
qiyuezhibo@qiyuezhibo-virtual-machine:~$ sudo cp /home/qiyuezhibo/Desktop/pathtoprocsched/procsched.h /usr/include
[sudo] password for qiyuezhibo:
qiyuezhibo@qiyuezhibo-virtual-machine:~$ gcc -shared -fPIC procsched.c -o libprocsched.so
gcc: error: procsched.c: No such file or directory
gcc: fatal error: no input files
compilation terminated.
qiyuezhibo@qiyuezhibo-virtual-machine:~$ cd /home/qiyuezhibo/Desktop/
qiyuezhibo@qiyuezhibo-virtual-machine:~/Desktop$ gcc -shared -fPIC procsched.c -o libprocsched.so
qiyuezhibo@qiyuezhibo-virtual-machine:~/Desktop$ sudo cp /home/qiyuezhibo/Desktop/libprocsched.so /usr/lib
qiyuezhibo@qiyuezhibo-virtual-machine:~/Desktop$
```

Last Step!

Create check.c file to check:

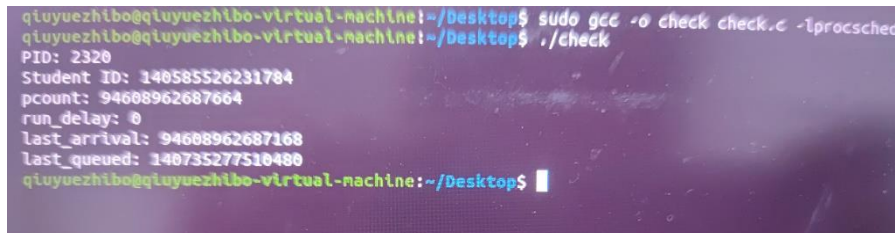
```
#include <procsched.h>
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdint.h>

int main() {
    pid_t mypid = getpid();
    printf("PID: %d\n", mypid);
    struct proc_segs info;

    if (procsched(mypid, &info) == 0) {
        printf("Student ID: %lu \n", info.mssv);
        printf("pcount: %lu \n", info.pcount);
        printf("run_delay: %llu \n", info.run_delay);
        printf("last_arrival: %llu\n", info.last_arrival);
        printf("last_queued: %llu\n", info.last_queued);
    } else {
        printf("Cannot get information from the process %d\n", mypid);
    }
    // If necessary, uncomment the following line to make this program run
    // long enough so that we could check out its maps file
    // sleep(100);
}
```

sudo gcc -o check check.c -lprocsched

./check



```
qiuyuezhibo@qiuyuezhibo-virtual-machine:~/Desktop$ sudo gcc -o check check.c -lprocsched
qiuyuezhibo@qiuyuezhibo-virtual-machine:~/Desktop$ ./check
PID: 2320
Student ID: 140585526231784
pcount: 94608962687664
run_delay: 0
last_arrival: 94608962687168
last_queued: 140735277510480
qiuyuezhibo@qiuyuezhibo-virtual-machine:~/Desktop$
```

Final output:

PID: 2320

Student ID: 140585526231784

pcount: 94608962687664

run_delay: 0

last_arrival: 94608962687168

last_queued: 140735277510480

Succeed and finish!

If the compilation ends successfully, copy the output file to `/usr/lib`. (Remember to add `sudo` before `cp` command).

QUESTION: Why we must put `-share` and `-fpic` option into `gcc` command?.

-Share: indicates that a shared library is generated.

-FPIC: indicates that using position independent cache.

The shared object can be loaded to different positions by different processes.

The purpose is to fulfill the request that compile source code as a shared to allow user to integrate our system call to their applications